

"unsupervised learning"

7

- unsupervised learning

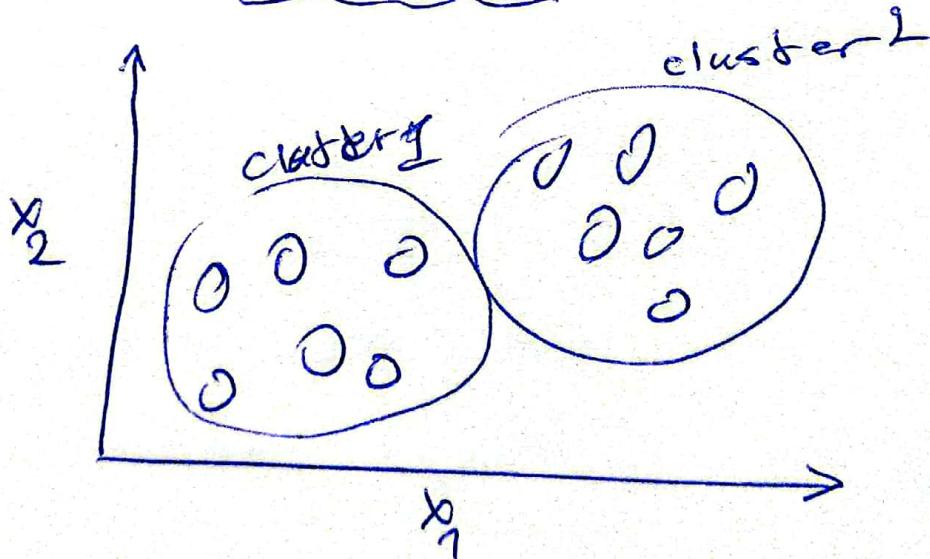
- clustering
- Anatomy detection

- Recommender Systems

- Reinforcement learning

other

clustering



Training sets $\{x^{(1)}, x^{(2)}, \dots, x^{(m)}\}$

- there is no target (label) for features.

(2)

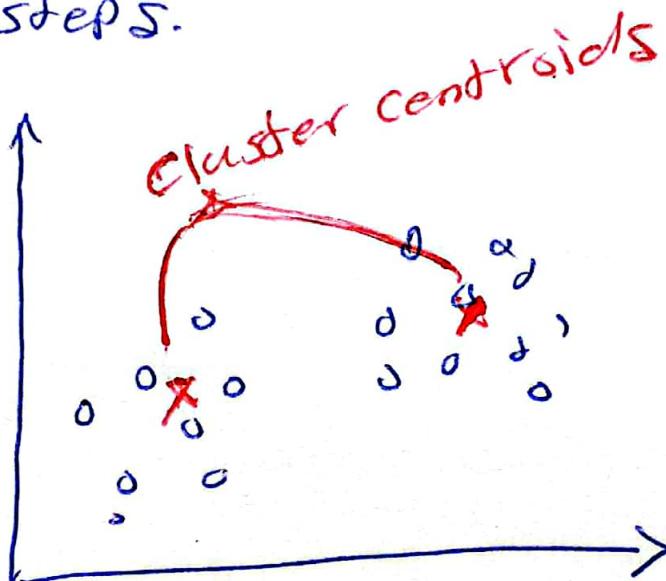
K-means algorithm (a clustering algorithm)

init: K-centroids (randomly choose)

Step 1: Assign each point to its closest centroid.

Step 2: Recompute the centroids.

repeat the steps 5.



Steps in detail

1. Randomly initialize K cluster centroids

$\vec{u}_1, \vec{u}_2, \dots, \vec{u}_K$ (u_i has the same dimensions like a training set)
each example

2. Repeat {

A. Assign points to cluster centroids

for $i = 1$ to m

$c^{(i)} := \text{index}(\text{from } 1 \text{ to } K) \text{ of cluster centroid}$
closest to $x^{(i)}$.

$$\min_k \|x^{(i)} - \mu_k\|^2$$

B. move cluster centroids

for $k = 1 \dots K$

μ_k = average (mean) of points assigned to cluster

ex.
 {
 }

$$\text{exp} \rightarrow \mu_k = \frac{1}{4} [x^{(1)} + x^{(5)} + x^{(6)} + x^{(10)}]$$

centroid

if during the process one cluster (K_i)

has no data point assigned to it

one should eliminate that cluster centroid.

$$K = K - 1$$

R-means optimization objective.

$c^{(i)}$ index of cluster ($1, 2, \dots, k$) to which example $x^{(i)}$ is currently assigned

μ_k cluster centroid k

$\mu_c^{(i)}$ cluster centroid of cluster to which example $x^{(i)}$ has been assigned.

(q)

cost function (K-means) \rightarrow

$$j(C^{(1)}, \dots, C^{(m)}, \mu_1, \dots, \mu_k) = \frac{1}{m} \sum_{i=1}^m \|x^{(i)} - \mu_{c(i)}\|^2$$

min

$$C^{(1)}, \dots, C^{(m)} \ j(C^{(1)}, \dots, C^{(m)}, \mu_1, \dots, \mu_k)$$
$$\mu_1, \dots, \mu_k$$

- K-means algorithm is trying to ~~find minimum~~
~~the~~ find assignments of points of ~~cluster~~
clusters centroid as well as find locations of
clusters centroids that minimizes the
squared distance.

- cost function of K-means called
Distortion function

initializing K-means

K-means algorithm

step 0: randomly initialize ~~K~~ ^{clusters} centroids $\mu_1, \mu_2, \dots, \mu_k$

Repeat {

 Step 1: Assign points to cluster centroids

 Step 2: move cluster centroids

Step 0

choose $k < m$

number of training examples
→ (5)

- randomly pick K training examples.
- set u_1, u_2, \dots, u_k equal to those K examples.
- ran the algorithm multiple times and each time with different set of randomly chosen cluster centroids is a good way to find ~~the~~ a better local minimum for cost function (distortion function)

for $i = 1$ to 100 { 50 to 1000 times is common

Randomly initialize K-means.

Run K-means. Get $C^{(1)}, \dots, C^{(m)}, u_1, u_2, \dots, u_k$
compute cost function (distortion)

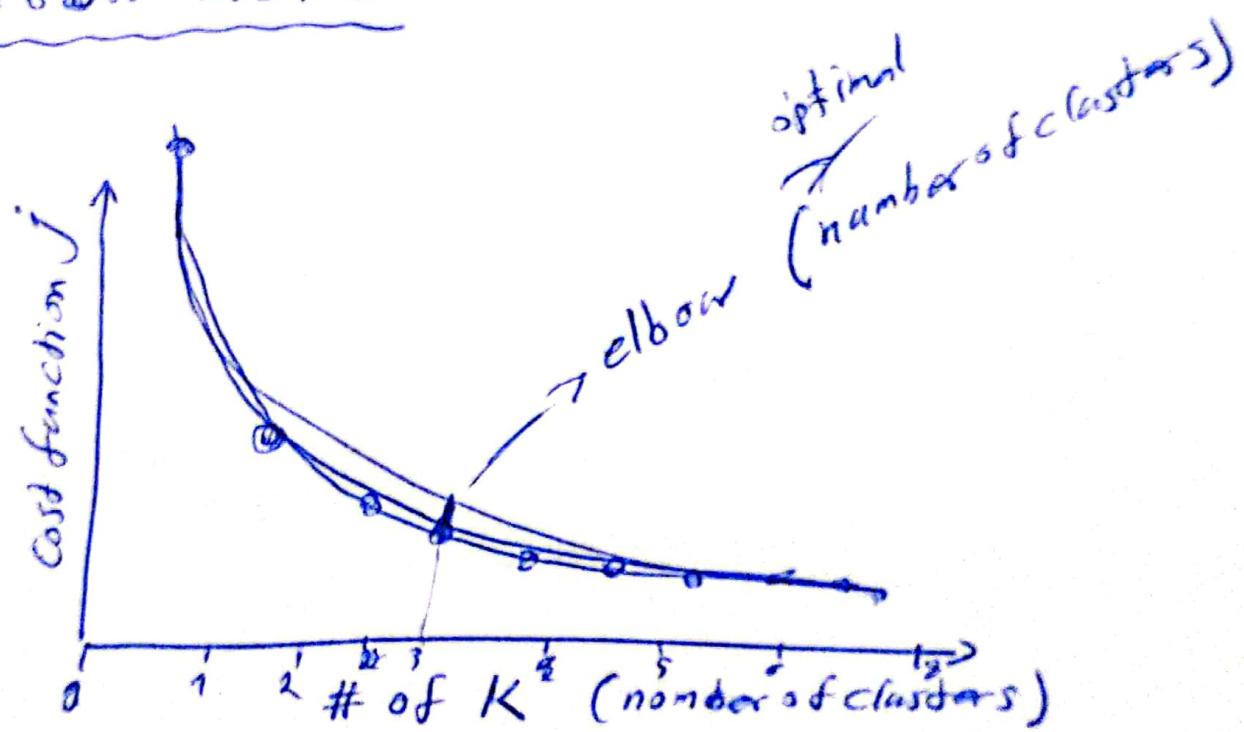
$$J(C^{(1)}, \dots, C^{(m)}, u_1, u_2, \dots, u_k)$$

PICK set of clusters that gave lowest cost J

choosing the number of clusters (# K ?) ⑥

choosing the value of K

Elbow method



problem, more K (cluster centroids) leads to to reduce cost function normally, so with elbow method in the most cases number of clusters are more than necessary.

choosing ~~the~~ the value of K

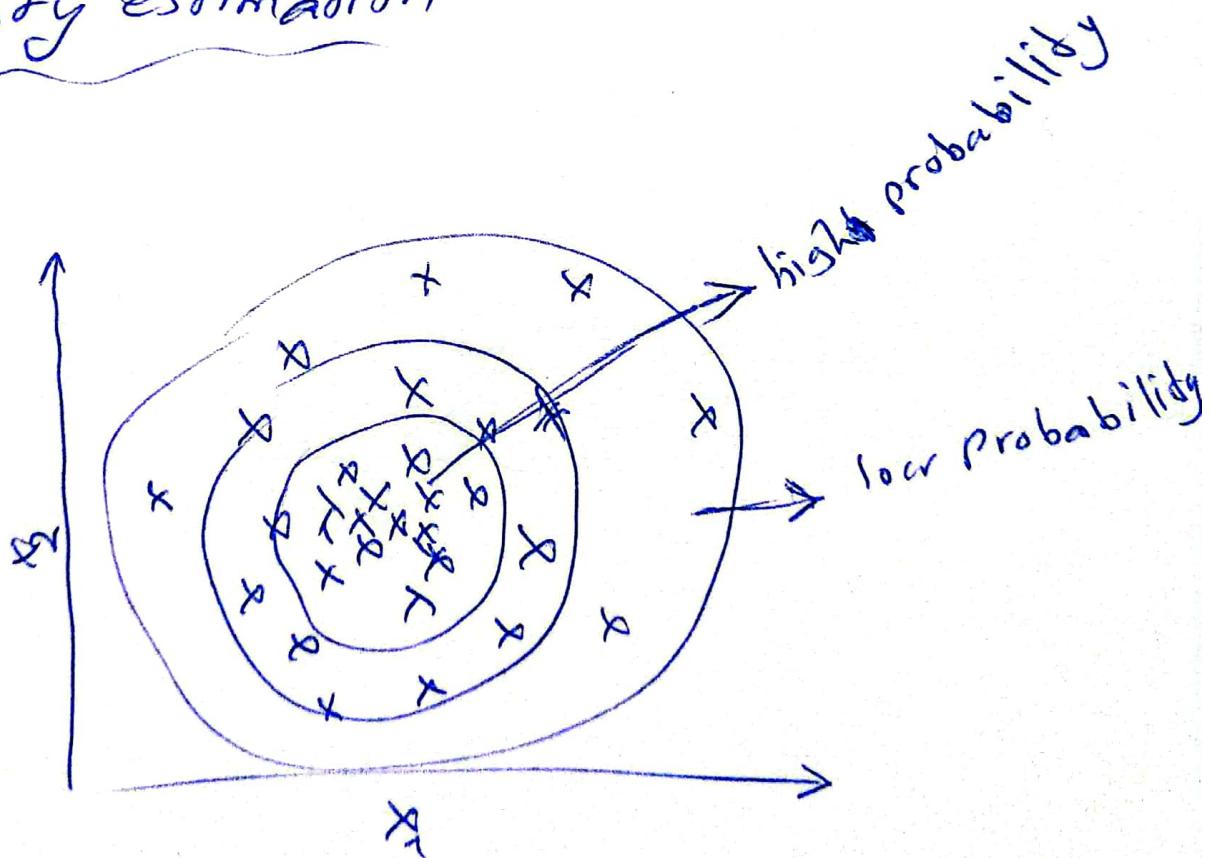
(7)

often, you want to get clusters for some later (downstream) purpose.

evaluate K-means based on how well it performs on that later purpose.

"Anomaly detection"

Density estimation



Dataset - $\{x^{(1)}, x^{(2)}, \dots, x^{(m)}\}$ epsilon

model $P(x)$ = probability of x being seen in dataset
is x_{test} anomalous? Yes if $P(x_{test}) < \epsilon$

Anomaly detection example:

- Fraud detection:

- manufacturing:

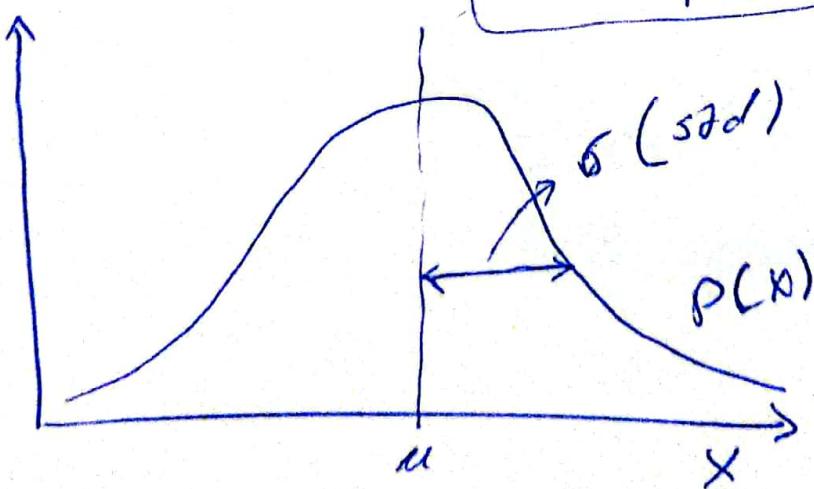
- monitoring computers in a data center.

Gaussian distribution

Say x is a number

- probability of x is determined by a Gaussian with mean μ and Variance σ^2 .

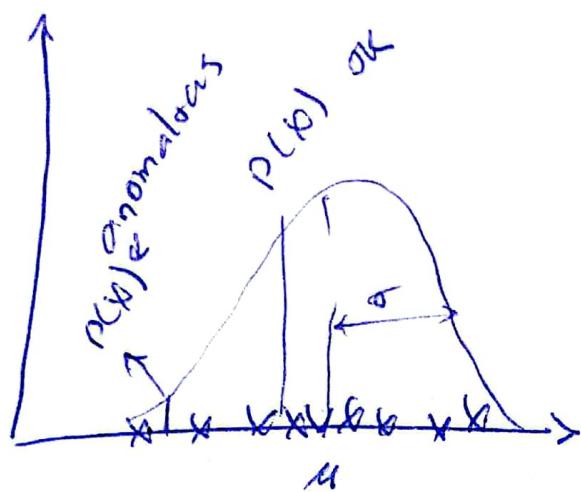
$$P(x) = \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{(x-\mu)^2}{2\sigma^2}}$$



Parameter estimation

9

Data set: $\{x^{(1)}, x^{(2)}, \dots, x^{(m)}\}$



$$\mu = \frac{1}{m} \sum_{i=1}^m x^{(i)}$$

maximum likelihood for
 μ and σ

$$\sigma^2 = \frac{1}{m} \sum_{i=1}^m (x^{(i)} - \mu)^2$$

sometimes

note: here ~~is just~~ the data set has just one feature.

Anomaly Detection Algorithm.

Training set: $\{\vec{x}^{(1)}, \vec{x}^{(2)}, \dots, \vec{x}^{(m)}\}$

(each example $\vec{x}^{(i)}$ has n features.)

Density estimation

$$P(\vec{x}) = P(x_1; \mu_1, \sigma_1^2) \times P(x_2; \mu_2, \sigma_2^2) \times P(x_3; \mu_3, \sigma_3^2) \times \dots \times P(x_n; \mu_n, \sigma_n^2)$$

$$= \prod_{j=1}^n P(x_j; \mu_j, \sigma_j^2)$$

Algorithm

1. choose n features x_i that you think might be indicative of anomalous examples.

2. fit parameters $\mu_1, \dots, \mu_n, \sigma_1^2, \dots, \sigma_n^2$

$$\mu_j = \frac{1}{m} \sum_{i=1}^m x_j^{(i)}$$

$$\sigma_j^2 = \frac{1}{m} \sum_{i=1}^m (x_j^{(i)} - \mu_j)^2$$

vectorized formula

$$\vec{\mu} = \frac{1}{m} \sum_{i=1}^m \vec{x}^{(i)}$$

3. Given new example x , compute $P(x)$:

$$P(x) = \prod_{j=1}^n P(x_j; \mu_j, \sigma_j^2) = \prod_{j=1}^n \frac{1}{\sqrt{2\pi\sigma_j^2}} \exp\left(-\frac{(x_j - \mu_j)^2}{2\sigma_j^2}\right)$$

Anomaly if $P(x) < \epsilon$

"The importance of real-number evaluation"

- when developing a learning algorithm (choosing feature, etc.), making decisions is much easier if we have a way of evaluating our learning algorithm.

Algorithm evaluation (Anomaly detection)

- fit model $P(x)$ on training set $x^{(1)}, x^{(2)}, \dots, x^{(n)}$
- on a cross validation/test example x , predict

$$y = \begin{cases} 1 & \text{if } P(x) < \epsilon \text{ (anomaly)} \\ 0 & \text{if } P(x) \geq \epsilon \text{ (normal)} \end{cases}$$

- possible evaluation metrics:

- True positive, false positive, false negative, true negative.
- Precision / Recall
- F_1 -score

use cross validation to choose parameter $\underline{\epsilon}$.
 \hat{x}
 set

Anomaly detection VS. Supervised learning

(72)

Anomaly detection

very small number of positive examples ($g=1$) - (0-20 is common).

Large number of negative ($g=0$) examples.

Supervised learning

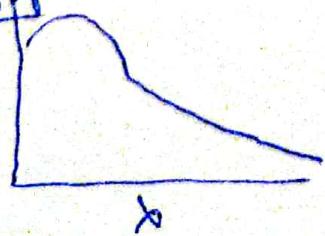
larger number of positive and negative examples.

choosing what features to use.

1. choose features which are normally distributed

if they are not normally distributed try to transform them. ~~anti~~

ex:



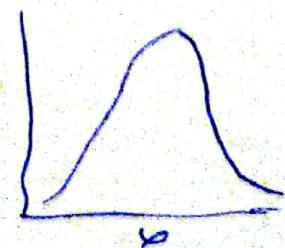
np.sqrt(x)

np.log(x)

or

np.log(x+1)

or $x = \frac{1}{x}$



Error analysis for anomaly detection

(73)

want $p(x) \geq \epsilon$ for normal examples x .
 $p(x) < \epsilon$ for anomalous examples.

most common problems:

$p(x) \approx \epsilon$ for both normal and anomalous examples.

"Recommender Systems"

an example: movie rating

$r(i, j) = 1$ if user j has rated movie i (0 otherwise)

$y(i, j)$ = rating given by user j on movie i (if defined)

$w^{(j)}, b^{(j)}$ = parameters for user j

$x^{(i)}$ = feature vector for movie i

for user j and movie i , predict rating: $w^{(j)} \cdot x^{(i)} + b^{(j)}$

$m(j)$ = number of movies rated by user j

To learn $w^{(j)}, b^{(j)}$

~~cost function~~

$$J(w^{(j)}, b^{(j)}) = \frac{1}{2m(j)} \sum_{i: r(i, j)=1} (w^{(j)} \cdot x^{(i)} + b^{(j)} - y^{(i, j)})^2 + \frac{1}{m(j)} \sum_k (w_k^{(j)})^2$$

cost function

$$J(w^{(1)}, b^{(1)}) = \frac{1}{2} \sum_{i: r(i,j)=1} (w_1^{(1)} \cdot x^{(i)} + b^{(1)} - y^{(i,j)})^2 + \frac{1}{2} \sum_{k=1}^n (w_k^{(1)})^2$$

To learn Parameters $w^{(1)}, b^{(1)}, w^{(2)}, b^{(2)}, \dots, w^{(n)}, b^{(n)}$ for
 numbers of users

all users :

$$\frac{1}{2} \sum_{j=1}^{n_u} \sum_{i: r(i,j)=1} (w^{(j)} \cdot x^{(i)} + b^{(j)} - y^{(i,j)})^2 + \frac{1}{2} \sum_{j=1}^{n_u} \sum_{k=1}^n (w_k^{(j)})^2$$

"collaborative filtering algorithm"

Given : $w^{(1)}, b^{(1)}, w^{(2)}, b^{(2)}, \dots, w^{(n)}, b^{(n)}$

To learn $x^{(1)}$:

$$J(x^{(1)}) = \frac{1}{2} \sum_{j: r(i,j)=1} (w^{(j)} \cdot x^{(i)} + b^{(j)} - y^{(i,j)})^2 + \frac{1}{2} \sum_{k=1}^n (x_k^{(i)})^2$$

minimizing the cost function as a function of $x^{(1)}$

To learn $x^{(1)}, x^{(2)}, \dots, x^{(n)}$:

$$\frac{1}{2} \sum_{i=1}^{n_m} \sum_{j: r(i,j)=1} (w^{(j)} \cdot x^{(i)} + b^{(j)} - y^{(i,j)})^2 + \frac{1}{2} \sum_{i=1}^{n_m} \sum_{k=1}^n (x_k^{(i)})^2$$

$$\Rightarrow J(x^{(1)}, x^{(2)}, \dots, x^{(n)})$$

in collaborative filtering algorithm we can learn

the features for a given ~~dataset~~ dataset.

(in contrast to most machine learning algorithms, ~~which~~ in which the features have to be externally given)

collaborative filtering algorithm

① cost function to learn $w^{(1)}, b^{(1)}, \dots, w^{(n_a)}, b^{(n_a)}$:

$$\min_{w^{(1)}, b^{(1)}, \dots, w^{(n_a)}, b^{(n_a)}} \frac{1}{2} \sum_{j=1}^{n_a} \sum_{i:r(i,j)=1} (w^{(i)} \cdot x^{(i)} + b^{(j)} - y^{(i,j)})^2 + \frac{1}{2} \sum_{j=1}^{n_a} \sum_{k=1}^n (w_k^{(j)})^2$$

② cost function to learn $x^{(1)}, \dots, x^{(n_m)}$:

$$\min_{x^{(1)}, \dots, x^{(n_m)}} \frac{1}{2} \sum_{i=1}^{n_m} \sum_{j:r(i,j)=1} (w^{(i)} \cdot x^{(i)} + b^{(j)} - y^{(i,j)})^2 + \frac{1}{2} \sum_{i=1}^{n_m} \sum_{k=1}^n (x_k^{(i)})^2$$

put them together (1 and 2):

$$\frac{1}{2} \sum_{(i,j):r(i,j)=1} (w^{(j)} \cdot x^{(i)} + b^{(j)} - y^{(i,j)})^2 + \frac{1}{2} \sum_{j=1}^{n_a} \sum_{k=1}^n (w_k^{(j)})^2 + \sum_{i=1}^{n_m} \sum_{k=1}^n (x_k^{(i)})^2.$$

overall cost function for learning

Params and features in collaborative filtering algorithm.

Gradient descent for collaborative filtering

10

Linear regression (Collaborative filtering)

repeat {

$$w_i^{(j)} = w_i^{(j)} - \alpha \frac{\partial}{\partial w_i^{(j)}} j(w, b, x)$$

$$b^{(j)} = b^{(j)} - \alpha \frac{\partial}{\partial b^{(j)}} j(w, b, x)$$

$$x_k^{(i)} = x_k^{(i)} - \alpha \frac{\partial}{\partial x_k^{(i)}} j(w, b, x)$$

g

Parameters : w, b, x

— — —

"Binary labels" (Collaborative filtering)

from regression to binary classification

regression, predict $y^{(i,j)}$ as $w^{(i)} \cdot x^{(i)} + b^{(j)}$

for binary labels,

predict the probability of $y^{(i,j)} = 1$

is given by $g(w^{(i)} \cdot x^{(i)} + b^{(j)})$

where $g(z) = \frac{1}{1+e^{-z}}$

in collaborative filtering algorithm we can learn

the features for a given ~~dataset~~ dataset.

(in contrast to most machine learning algorithms, ~~which~~ in which the features have to be externally given)

collaborative filtering algorithm

① cost function to learn $w^{(1)}, b^{(1)}, \dots, w^{(n_a)}, b^{(n_a)}$:

$$\min_{w^{(1)}, b^{(1)}, \dots, w^{(n_a)}, b^{(n_a)}} \frac{1}{2} \sum_{j=1}^{n_a} \sum_{i: r(i,j)=1} (w^{(1)} \cdot x^{(i)} + b^{(j)} - y^{(i,j)})^2 + \frac{1}{2} \sum_{j=1}^{n_a} \sum_{k=1}^n (w_k^{(j)})^2$$

② cost function to learn $x^{(1)}, \dots, x^{(n_m)}$:

$$\min_{x^{(1)}, \dots, x^{(n_m)}} \frac{1}{2} \sum_{i=1}^{n_m} \sum_{j: r(i,j)=1} (w^{(j)} \cdot x^{(i)} + b^{(j)} - y^{(i,j)})^2 + \frac{1}{2} \sum_{i=1}^{n_m} \sum_{k=1}^n (x_k^{(i)})^2$$

put them together (1 and 2):

$$\frac{1}{2} \sum_{(i,j): r(i,j)=1} (w^{(j)} \cdot x^{(i)} + b^{(j)} - y^{(i,j)})^2 + \frac{1}{2} \sum_{j=1}^{n_a} \sum_{k=1}^n (w_k^{(j)})^2 + \sum_{i=1}^{n_m} \sum_{k=1}^n (x_k^{(i)})^2$$

overall cost function for learning

Params and features in collaborative filtering algorithm.

Gradient descent for collaborative filtering

(10)

Linear regression (Collaborative filtering)

repeat {

$$w_i^{(j)} = w_i^{(j)} - \alpha \frac{\partial}{\partial w_i^{(j)}} j(w, b, x)$$

$$b^{(j)} = b^{(j)} - \alpha \frac{\partial}{\partial b^{(j)}} j(w, b, x)$$

$$x_k^{(i)} = x_k^{(i)} - \alpha \frac{\partial}{\partial x_k^{(i)}} j(w, b, x)$$

}

Parameters : w, b, x

— — —

"Binary labels" (Collaborative filtering)

from regression to binary classification

regression: predict $y^{(i,j)}$ as $w^{(j)} \cdot x^{(i)} + b^{(j)}$

for binary labels:

predict the probability of $y^{(i,j)} = 1$

is given by $g(w^{(j)} \cdot x^{(i)} + b^{(j)})$

$$\text{where } g(z) = \frac{1}{1+e^{-z}}$$

cost function for binary application
(collaborative filtering)

$$y^{(i,j)}; f_{(w,b,x)}(x) = g(w^T \cdot x^{(i)} + b^{(j)})$$

 cost function:

$$J(w, b, x) = \sum_{(i,j) : r(i,j)=1} L(f_{(w,b,x)}(x), y^{(i,j)})$$

$$\hookrightarrow L(f_{(w,b,x)}(x), y^{(i,j)}) = -y^{(i,j)} \log(f_{(w,b,x)}(x)) - (1-y^{(i,j)}) \log(1 - f_{(w,b,x)}(x))$$

~~recommendation~~

mean normalization

$$\begin{array}{l}
 \text{user} \\
 \text{movie} \\
 \begin{bmatrix}
 5 & 5 & 0 & 0 & ? \\
 5 & ? & 2 & 0 & ? \\
 ? & 4 & 0 & ? & ? \\
 0 & 0 & 5 & 0 & 3
 \end{bmatrix}
 \begin{bmatrix}
 2.5 \\
 2.5 \\
 2 \\
 2.25
 \end{bmatrix}
 \ominus \mu = \begin{bmatrix}
 2.5 \\
 2.5 \\
 2 \\
 2.25
 \end{bmatrix} = \begin{bmatrix}
 2.5 & 2.5 & 2.5 \\
 \hline
 & &
 \end{bmatrix} \\
 \qquad \qquad \qquad y^{(i,j)}
 \end{array}$$

Tensorflow implementation of collaborative filtering

18

Custom Training loop

$$f(x) = g$$
$$J = (\hat{w}x - 1)^2$$

gradient descent algorithm

$$w = w - \alpha \frac{d}{dw} J(w, b)$$

$f(w, b) = 0$ for this example

initialize w with 3

using tensorflow

$$w = tf.Variable(3.0)$$

$$x = 1.0$$

$$y = 1.0 \# \text{target value}$$

$$\alpha = 0.07$$

$$\text{iterations} = 30$$

for iter in range(iterations):

Use Tensorflow's gradient tape to record the steps.

Used to compute the cost J , to enable auto differentiation.

with `tf.GradientTape()` as tapes:

$$fw_b = w \times x$$

$$\text{cost}_j = (fw_b - y)^2$$

Use the GradientTape to calculate the gradients of the cost with respect to the parameter w .

$$[\frac{dJ}{dw}] = \text{tape.gradient}(\text{cost}_j, [w])$$

ran one step of GD by updating the value of wbo 79

to reduce the cost

w = assign_sub(-alpha * djdw)

Implementation in Tensorflow (collaborative filtering)

optimizer = keras.optimizers.Adam(learning_rate=1e-1)

iters = 200

for iter in range(iters):

with tf.GradientTape() as tape:

cost_values = cost.CostFunc(x, ws, b, Ynorm, R, numvars,
num_ratings, lambdab)

grads = tape.gradient(cost_value, [x, ws, b])

optimizer.apply_gradients(zip(grads, [x, ws, b]))

Limitations of Collaborative Filtering

- Cold start problem: How do

rank near items that few users have rated?

• show something reasonable to new users ~~new~~ who have rated few items?

- use side information about items or users;

• item: genre, moviestars, studios, ...

- user Demographics (age, gender, location), expressed preferences, ...

20

collaborative filtering vs. content-based filtering.

Collaborative filtering: Recommend items to you based on ratings of users who gives similar ratings as you.

content-based filtering: Recommend items to you based on features of user and item to find good match.

examples of user and item features

user features:

[Age
Gender
Country
movies watched
↳ average rating per genre]

} $x_u^{(j)}$ for user j

movie features:

[Year
Genre/Genres
Reviews
Average rating]

} $x_m^{(i)}$ for movie i

Predict rating of user j on movie i as

(27)

~~intend~~
collaborative filtering

$$\boxed{w^{(j)} \cdot x^{(i)} + b^{(j)}}$$
$$\boxed{v_u^{(j)} \cdot v_m^{(i)}} \quad x \leftarrow \text{content based filtering}$$

computed from $x_u^{(j)}$ computed from $x_m^{(i)}$

exp:

user's preferences $v_u^{(j)}$ Genres
 $\begin{bmatrix} 4 \\ 0.9 \\ 0.7 \\ ; \\ ; \\ 0.0 \end{bmatrix}$ how much user j likes each genre

both have to have the same dimension

movie features $v_m^{(i)}$ romance how much end movie i is action or -
 $\begin{bmatrix} 0.5 \\ 0.2 \\ ; \\ 3.5 \end{bmatrix}$ action

deep learning for content-based filtering
neural network architecture

$x_u \rightarrow v_u$ user network

$x_m \rightarrow v_m$ movie network

User network

$$(x_a \rightarrow v_a)$$



Movie network

$$(x_m \rightarrow v_m)$$



Same output layer size.

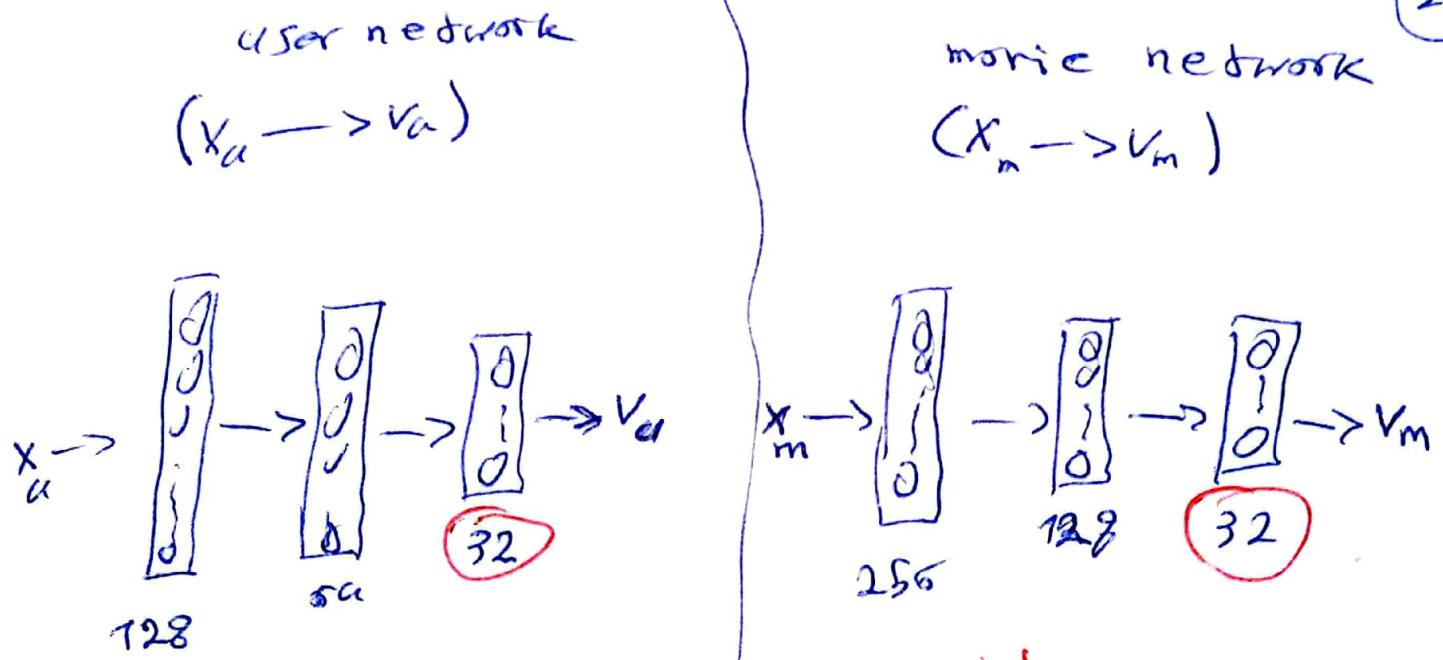
$$\text{Predict: } v_a^{(i)} \cdot v_m^{(j)}$$

dot product

$g(v_a^{(i)} \cdot v_m^{(j)})$ to predict the probability that $y^{(i,j)}$ is 1.

neural network architecture



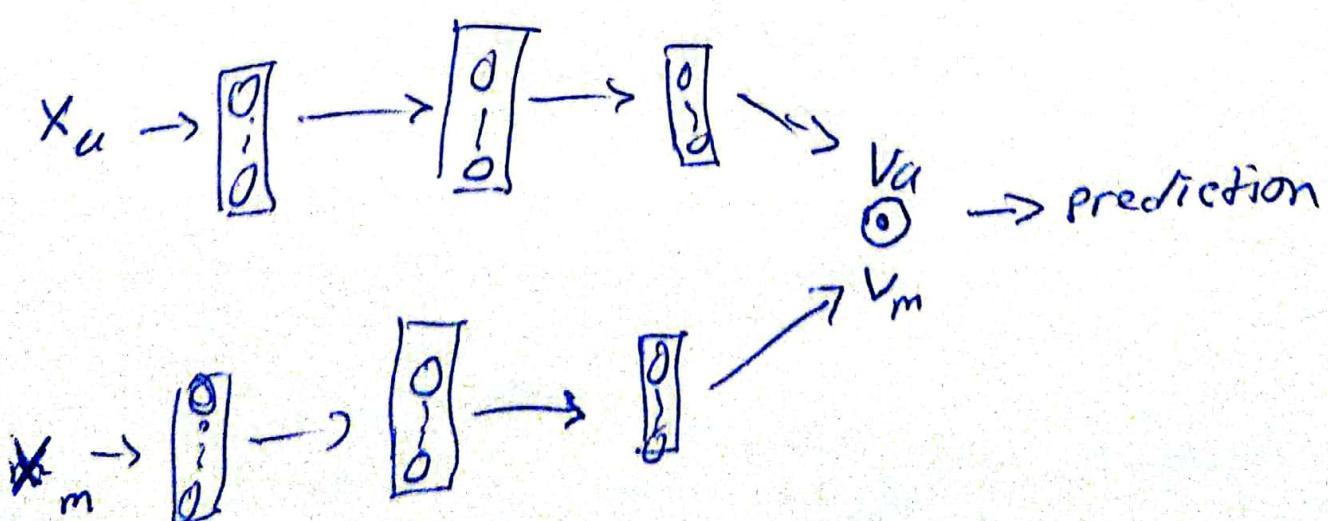


Same output layer size.

Predict: $v_a^{(i)} \cdot v_m^{(j)}$
dot product

or
 $g(v_a^{(i)} \cdot v_m^{(j)})$ to predict the probability that $g(i,j)$ is 1

neural network architecture



cost function for content-based filtering (23)

$$J = \sum_{(i,j) : r(i,j) > 0} (v_u^{(j)} \cdot v_m^{(i)} - y(i,j))^2 + \text{NN regularization term}$$

learned user and item vectors

$v_u^{(j)}$ is a vector of length 32 that describes user j with features $x_u^{(j)}$

$v_m^{(i)}$ is a vector of length 32 that describes movie i with features $x_m^{(i)}$

- To find movies similar to movie i :

$$\|v_m^{(k)} - v_m^{(i)}\|^2 \quad \begin{cases} \text{small} \\ \rightarrow \text{compute distance} \end{cases}$$

Note: this can be pre-computed ahead of time.

How to efficiently find recommendation from a large set of items?

Two steps; Retrieval & Ranking

(24)

Retrieval:

- generate large list of plausible item candidates

- e.g. ~100s
- retrieved items
 - 1) for each of the last 10 movies watched by the user, find 10 most similar movies.
 - 2) for most viewed 3 genres, find the top 10 movies
 - 3) Top 20 movies in the country

- combine retrieved items into list, removing duplicates and items already watched/purchased.

Ranking:

- take ~~the~~ retrieved list and rank using learned model

$$x_u \rightarrow [] \rightarrow [] \rightarrow [v_u]$$

$$x_m \rightarrow [] \rightarrow [v_m]$$

Predictions

- display ranked items to user.

Tensor flow implementation of content-based filtering. (25)

x_u
↓ input

user-MN = tf.keras.models.Sequential([

tf.keras.layers.Dense(256, activation="relu"),

tf.keras.layers.Dense(128, activation="relu"),

tf.keras.layers.Dense(32)

]) $\rightarrow v_u$ outPart

x_m
↓ input

item-MN = tf.keras.models.Sequential([

tf.keras.layers.Dense(256, activation="relu"),

tf.keras.layers.Dense(128, activation="relu"),

tf.keras.layers.Dense(32)

]) $\rightarrow v_m$ outPart

create the user input and point to the base network

input_user = tf.keras.layers.Input(shape=(num_user_features))

~~item-MN(input)~~

~~v_u = user-MN(input_user)~~

$v_u = \text{tf.math.l2_normalize}(v_u, axis=1)$

```

# create the item input and point to the base network. (25)
input_item = tf.keras.layers.Input(shape=(num_item_features))
vm = item_nn(input_item)
vm = tf.linalg.l2_normalize(vm, axis=1)

# measure the similarity of the two vector outputs
out_put = tf.keras.layers.Dot(axis=1)([vu, vm])

# specify the inputs and output of the model.
model = Model([input_user, input_item], out_put)

# specify the cost function
cost_fn = tf.keras.losses.MeanSquaredError()

```

Reducing the number of features

(PCA)
 Principal component analysis (an supervised learning
 algorithm to reduce the number of features)

PCA: finds new axis and coordinates

for example: a car dataset with two features

x_1 = length of car
 x_2 = height of car

can be combined and make
 a feature (size). ~~size~~
 & near

(27)

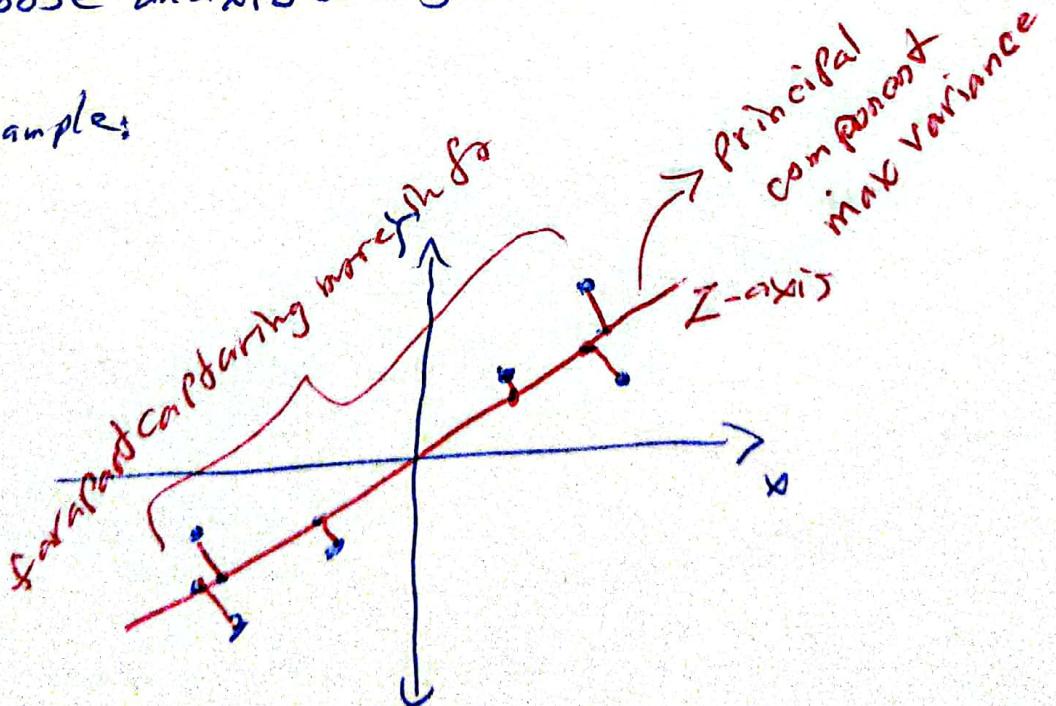
PCA: takes a data set with too many features (for example 1000) and compresses them into just 2 features which one can visualize in a 2 dimensional ~~graph~~ plot and gains information about all the data set (features) in just a visualization.

PCA Algorithm

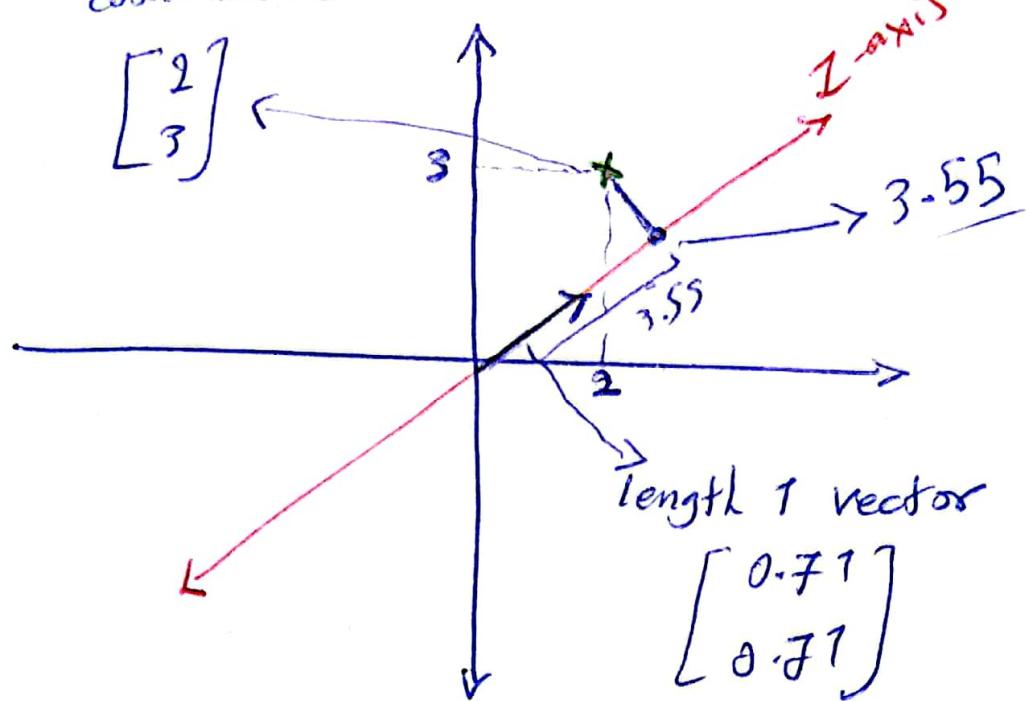
feature scaling

1. preprocessing: normalized the features to have the same mean (0) \rightarrow Z-score normalization.
 (Z)
2. Choose axis to project ~~the~~ all points (values).

examples:



coordinates



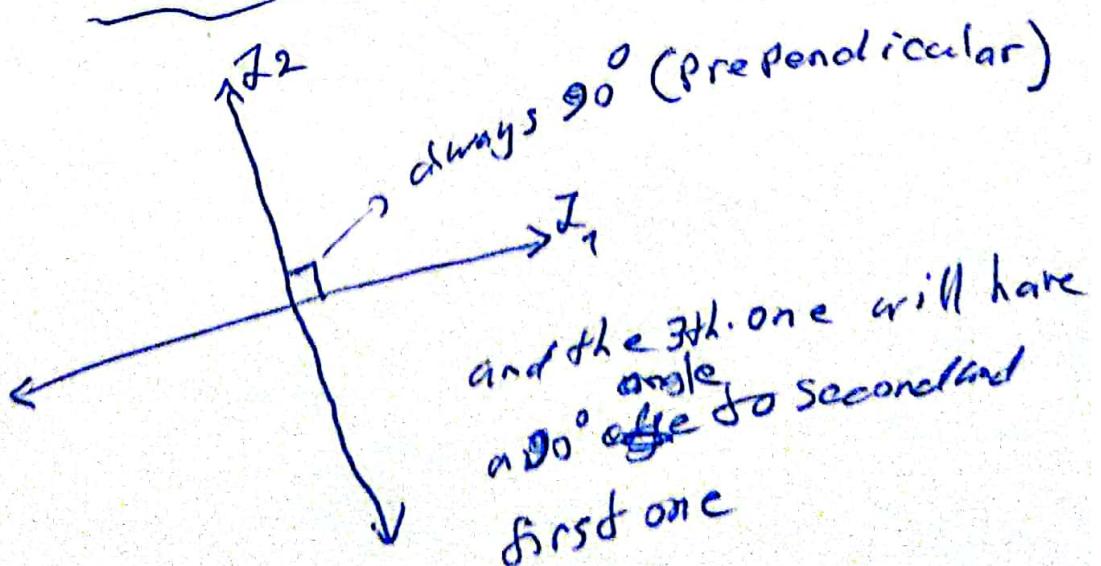
How to project \vec{x} with two dimensions in one dimension
given-1) a point

dot Product

$$\begin{bmatrix} 2 \\ 3 \end{bmatrix} \cdot \begin{bmatrix} 0.71 \\ 0.71 \end{bmatrix} = 2 \times 0.71 + 3 \times 0.71 = 3.55$$

projected point

more principal components



approximation to the original data

given $Z = 3.55 \rightarrow \text{as}$

find original $(X_1, X_2) = ?$

"reconstruction" \rightarrow approximation

$$3.55 \times \begin{bmatrix} 0.77 \\ 0.71 \end{bmatrix} = \begin{bmatrix} 2.52 \\ 2.52 \end{bmatrix}$$



approximate the original
Point

~~sklearn~~
(PCA in code) \rightarrow scikit-learn

1. scaled data: (feature scaling) Z-normalization

\rightarrow or none

2. "fit" the data to obtain 2(or 3) new axes (principal components)

performs
 \rightarrow fit ~~contains~~ mean normalization automatically, so on

No need to normalize data before "fit"

3. optionally examine how much variance is explained
by each principal component.

q. transform (project) the data onto the new axes. (30)

code

1. (optional) normalizations:

$X = np.array([[-1, 1], [2, 1], [3, 2], [-1, -1], [-2, -1], [-3, -2]])$

2. "fit"

$PCA_1 = PCA(n_components=1)$

$PCA_1.fit(X)$

3. variance explained

$PCA_1.explained_variance_ratio_$

0.992



In this example

4. projection (transform):

$X_{\text{trans-1}} = PCA_1.transform(X)$

$X_{\text{reduced-1}} = \cancel{PCA_1} PCA_1.inverse_transform(X_{\text{trans-1}})$

$\left(\begin{array}{l} array([[-1, 1], \\ [2, 1], \\ [3, 2], \\ [-1, -1], \\ [-2, -1], \\ [-3, -2]]) \end{array} \right)$

Applications of PCA

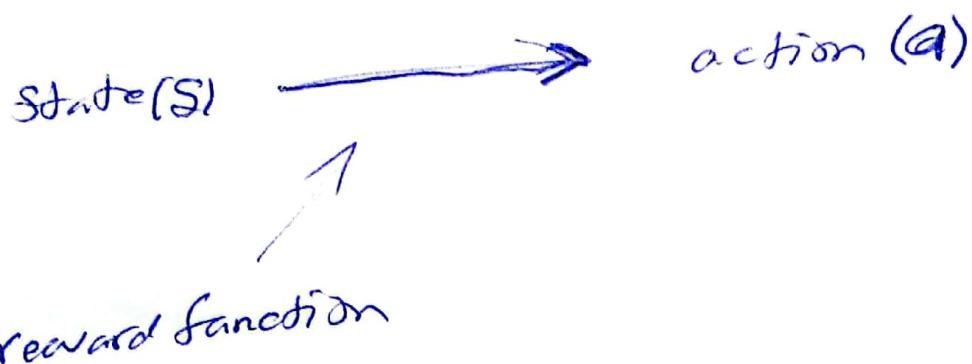
- appropriate for visualization purposes (reduce data to 2 or 3 Ds)
- less frequently used for:
 - Data compression (to reduce storage or transmission costs)
 - speeding up training of a supervised learning model.

Reinforcement learning

32

exps:

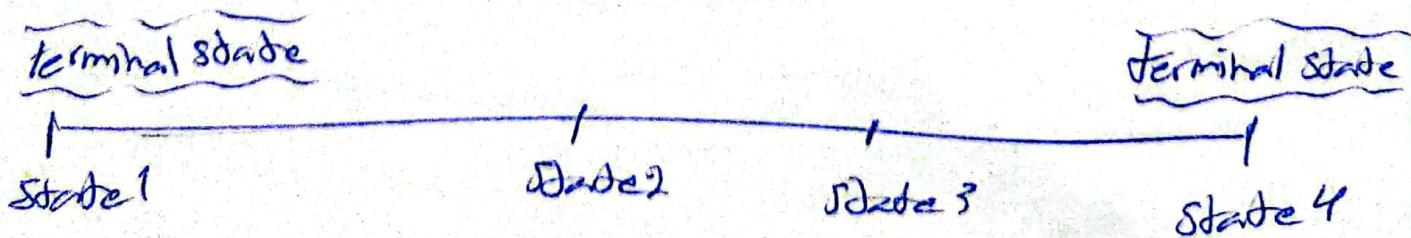
Position of helicopter → how to move control sticks



- positive reward: helicopter flying well $+1 \uparrow$
- negative reward, helicopter flying poorly $-1 \downarrow$

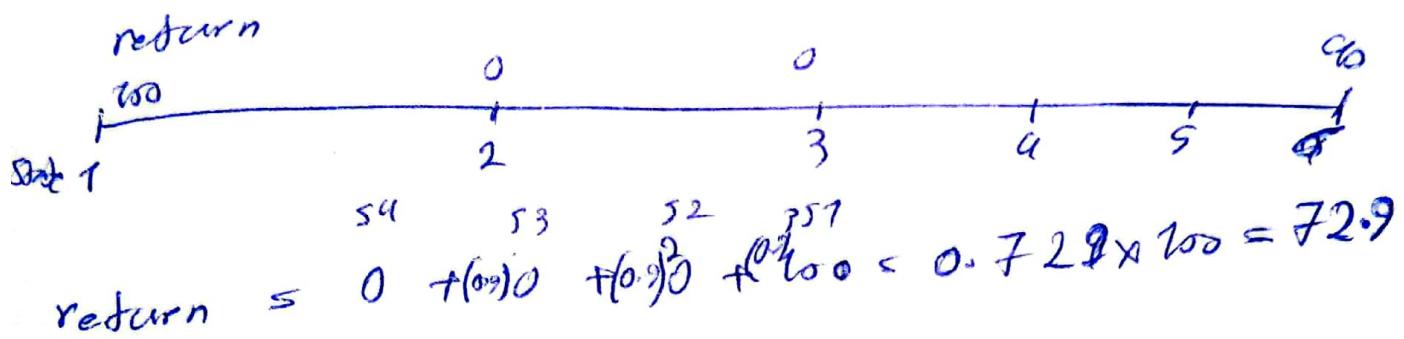
Applications of reinforcement algorithms

- controlling robots
- factory optimization
- financial (stock) trading
- Playing games (including video games)



(33)

return in reinforcement algorithms learning



$$\text{Return} = R_1 + \gamma R_2 + \gamma^2 R_3 + \gamma^3 R_4 + \dots$$

discount factor $\gamma = 0.9 \rightarrow$ In this example

a common choices of γ : 0.9, 0.99, 0.999, ...

a number near to 1

exp:

$$\boxed{\gamma = 0.5}$$

$$\text{return} = 0 + (0.5)0 + (0.5)^2 0 + (0.5)^3 100 = 12.5$$

making decision in reinforcement learning

state

s

action

$\pi \rightarrow a$

policy

find a policy π that tells good what action ($a = \pi(s)$) to take in every state (s) so as to maximize the return.

key concepts in reinforcement learning

example (mars rover) 34

state

6 states

actions

← →

rewards

100, 0.0, 0, 0, 0

discount factor
 γ

0.5

return

$$R_1 + \gamma R_2 + \gamma^2 R_3 + \gamma^3 R_4 + \dots$$

Policy

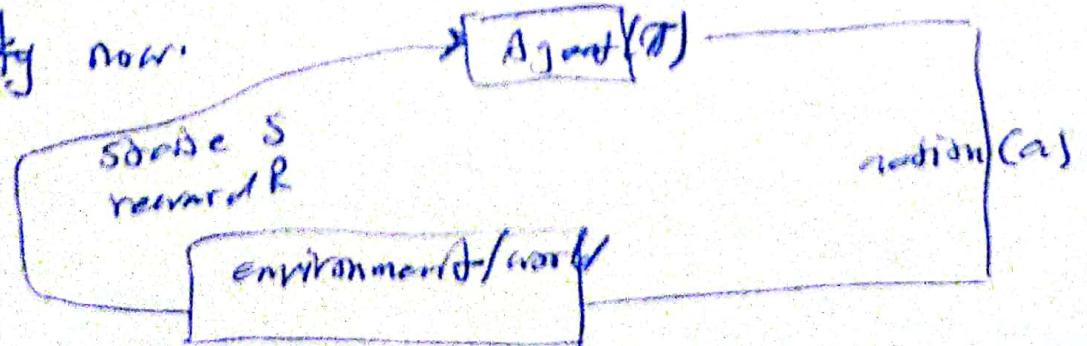
100	0	0	0	0	0
-----	---	---	---	---	---

← ← ← →

~~state 1 is 0~~

Mars Rover Decision Process (MDP)

-In a markov decision process the current step depends just on the last step. In other words the feature step depends on where you are currently now.



state-action value function

35

$Q(s, a)$ = Return if you

- start in state s .
- take action a (once)
- then behave optimally after that.

exp:

	100	50	25	12.5	20	90	return
	100	0	0	0	0	40	action
	1	2	3	4	5	6	reward

$$Q(2, \rightarrow) = 12.5$$

$$\leftarrow 0 + (0.5)0 + (0.5)^20 + (0.5)^3100$$

↓
Start in state $s(2)$

take action a (once)
(go to state 3)

behave optimally
(go again to state 2)

finally go to
state 1

The best possible return from state s is $\max_a Q(s, a)$ (30)

Bellman Equation

$Q(s, a)$ = return if you

- start in state s
- take action a (once)
- then behave optimally after that.

s : current state $R(s)$ = reward of current state

a : current action

s' : state you get to after taking action a .

a' : action that you take in state s'

Bellman equation

$$Q(s, a) = \underbrace{R(s)}_{\text{reward you get right away}} + \gamma \max_{a'} Q(s', a')$$

return from
behaving optimally
Starting from state
 s'

exp,

1	2	3	4	5	6	$s=2$
100	50	25	12.5	6.25	3.125	$a \rightarrow$
100	0	0	0	0	0	$s' = 3$

$$Q(2, \rightarrow) = R(2) + 0.5 \max_{a'} Q(3, a')$$

$$= 0 + (0.5 \times 25) < 12.5$$

$$Q(s,a) = R(s) + \gamma \max_{a'} Q(s', a')$$

$$R_1 + \gamma R_2 + \gamma^2 R_3 + \gamma^3 R_4 + \dots$$

$$P = R_1 + \gamma [R_2 + \gamma R_3 + \gamma^2 R_4 + \dots]$$

$Q(s, a)$

both have

the same result.

Random (Stochastic) environment

expected

Stochastic Reinforcement learning algorithm

$$\begin{aligned} \text{Return} &= \text{Average } (R_1 + \gamma R_2 + \gamma^2 R_3 + \gamma^3 R_4 + \dots) \\ &\leq E[R_1 + \gamma R_2 + \gamma^2 R_3 + \gamma^3 R_4 + \dots] \end{aligned}$$

Goal of reinforcement learning:

choose a policy $\pi(s) = a$ that will tell us what action a to take in state s so as to maximize the expected return

Bellman
Equation

$$Q(s, a) = R(s) + \gamma E[\max_{a'} Q(s', a')]$$

continuous State spaces

(38)

expts

coordinates of state: a car

$$S_s = \begin{bmatrix} x \\ y \\ \theta \\ \dot{x} \\ \dot{y} \\ \dot{\theta} \end{bmatrix} \rightarrow \begin{array}{l} x - \text{Position} \\ y - \text{Position} \\ \theta - \text{orientation} \\ \dot{x} - \text{velocity in X direction} \\ \dot{y} - \text{velocity in Y direction} \\ \dot{\theta} - \text{velocity of change in } \theta \end{array}$$

How quickly the angle of car changes.

exp 21 Autonomous Helicopter.

$$S = \begin{bmatrix} x \\ y \\ z \\ \phi \\ \theta \\ \omega \\ \dot{x} \\ \dot{y} \\ \dot{z} \\ \ddot{\phi} \\ \ddot{\theta} \\ \ddot{\omega} \end{bmatrix} \rightarrow \text{angles in 3D.}$$

("Lunar Lander")

actions:

f	- do nothing
	- left thruster
	- main thruster
	- right thruster

$$\mathcal{S} = \begin{bmatrix} x \\ y \\ \dot{x} \\ \dot{y} \\ \theta \\ \dot{\theta} \\ LG \\ r \end{bmatrix}$$

→ left and right
 → binary values

Reward Function

- Getting to landing path: 100 - 740
- Additional reward for moving toward/away from path
- crash: -100
- soft landing: +100
- leg grounded: +10
- fire main engines: -0.3
- fire side thruster: -0.03

Lunar Lander Problem

90

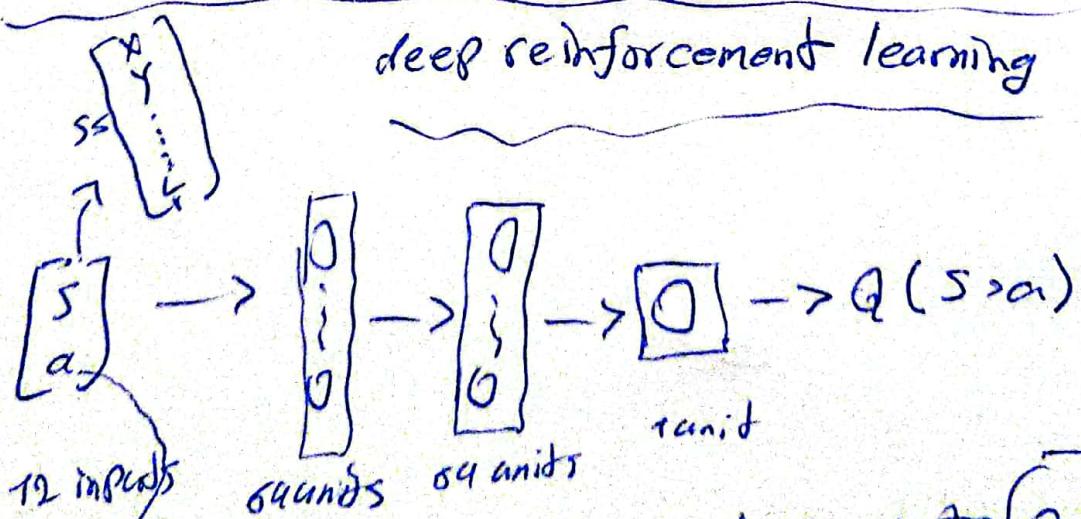
- Learn a policy π that, given

$$s \in \begin{bmatrix} x \\ y \\ \dot{x} \\ \dot{y} \\ \theta \\ \dot{\theta} \\ b \\ r \end{bmatrix}$$

- pick action $a = \pi(s)$ so as to maximize the return.

~~π~~.

- discount factor $\gamma = 0.985$



in a state s , use neural network to compute $(Q(s, \text{nothing}), Q(s, \text{left}), Q(s, \text{main}), Q(s, \text{right}))$
pick the action that maximizes $Q(s, a)$.

$$\begin{bmatrix} s \\ \alpha \end{bmatrix} = \begin{bmatrix} x \\ y \\ \vdots \\ i \\ Y \\ \bar{\alpha} \\ L \\ r \\ \alpha \\ 0 \\ 0 \\ 0 \\ 0 \end{bmatrix} \rightarrow Q(s, \text{nothing})$$

or

$$\begin{bmatrix} x \\ y \\ \alpha \\ \vdots \\ 0 \\ 1 \\ 0 \\ 0 \end{bmatrix} \rightarrow Q(s, \text{left})$$

or

$$\begin{bmatrix} x \\ y \\ \vdots \\ 0 \\ 0 \\ 1 \\ 0 \end{bmatrix} \rightarrow Q(s, \text{main})$$

or

$$\begin{bmatrix} x \\ y \\ \vdots \\ 0 \\ 0 \\ 1 \end{bmatrix} \rightarrow Q(s, \text{right})$$

$Q(s, a) \rightarrow$ output of ~~neural~~
 deep reinforcement
 network

y or target

Bellman equation

$$Q(s, a) = R(s) + \gamma \max_{a'} Q(s', a')$$

$\underbrace{x}_{\gamma} \quad \underbrace{\max_{a'}}_{y}$

create some labeled training examples,

$(s^{(1)}, a^{(1)}, R(s^{(1)}), s'^{(1)})$

$(s^{(2)}, a^{(2)}, R(s^{(2)}), s'^{(2)})$

\vdots

X	Y
$s^{(1)}, a^{(1)}$	$R(s^{(1)})$

$$R(s^{(1)}) + \gamma \max_{a'} Q(s'^{(1)}, a')$$

Learning algorithm (reinforcement)

↳ Deep Q network

(43)

- initialize neural network randomly as guess of $Q(s, a)$,
- Repeat {

Take actions in the lunar lander. Get $(s, a, r(s))$'s

store 10,000 most recent $(s, a, r(s), s')$ tuples-

↑
Replay buffer

Train neural network:

- create training set of 10000 examples using

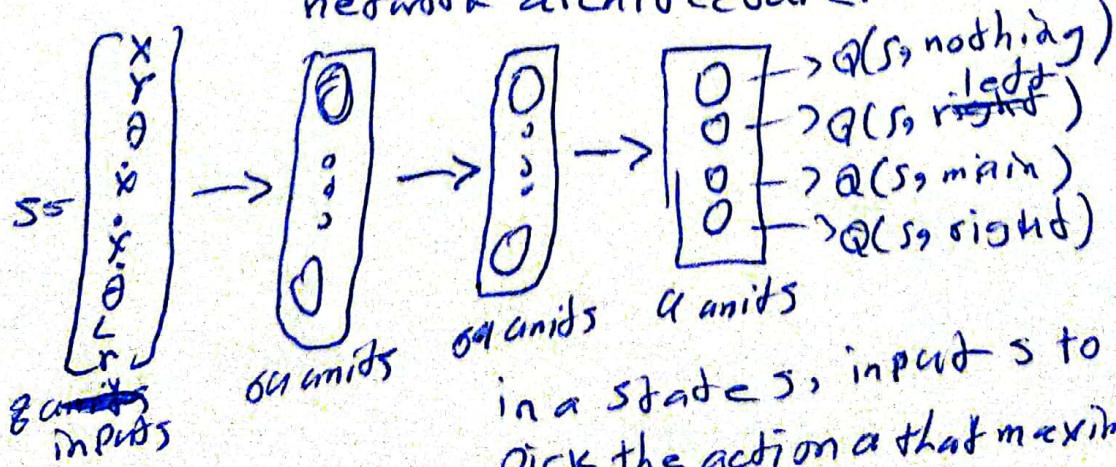
$$x = (s, a) \text{ and } y = r(s) + \gamma \max_{a'} Q(s', a')$$

- train Q_{new} such that $Q_{\text{new}}(s, a) \approx y$.

Set $Q = Q_{\text{new}}$

}

Algorithm (DQN) refinement: improved neural network architecture:



in a state s , input s to neural network
pick the action a that maximizes $Q(s, a)$

Algorithm reinforcement: ϵ -greedy policy

(44)

How to choose actions while still learning?

- in some state s

→ or "exploitation"

option 1:

- pick the action a that maximizes $Q(s, a)$.

option 2:

- with probability 0.95, pick the action a that maximizes $Q(s, a)$.

→ "greedy action"

- with probability 0.05, pick an action a randomly.

→ "exploration step"

→ the better option

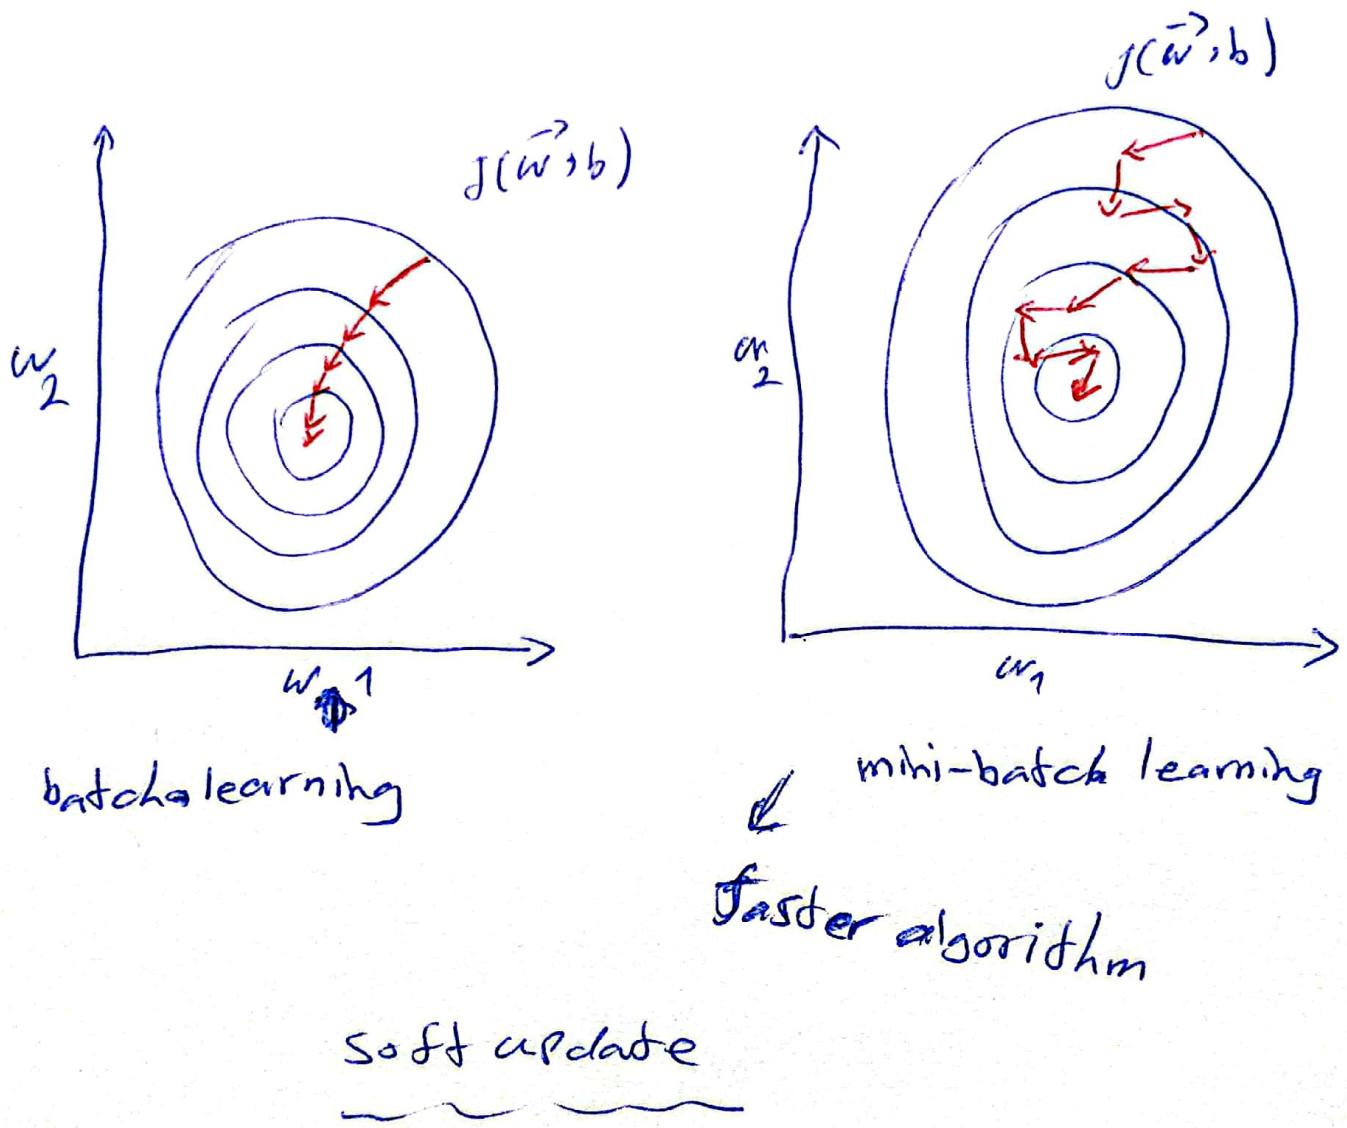
→ ϵ -greedy policy ($\epsilon = 0.05$)

start ϵ high
Gradually decrease

→ exp: start $\epsilon = 1$,

finally $\epsilon = 0.01$

Algorithm refinements: minibatch
and soft updates.



Set $Q = Q_{\text{new}}$

w, β

$w_{\text{new}}, \beta_{\text{new}}$

$$w = 0.07 w_{\text{new}} + 0.93 w$$

$$\beta = 0.07 \beta_{\text{new}} + 0.93 \beta$$

soft update.

Limitations of reinforcement learning.

(40)

- much easier to get to work in a simulation than a real robot!
- far fewer applications than supervised and unsupervised learning.
- But - exciting research direction with potential for feature applications.

~~The end of the machine-learning specialization~~

2024.9.27
i)

Lec 1